

Exercise 11 Raphael Michel and Florian Stoertz


July 14, 2017

```
In [39]: import matplotlib.pyplot as plt
         %matplotlib inline
         import numpy as np
         from sklearn.metrics import log_loss
```

1 Fisher information / Natural gradients

```
In [53]: Nsamples = 100
         dataAclass0 = np.random.multivariate_normal(np.array([1, 1]),
                                                    np.eye(2), Nsamples//2)
         dataAclass1 = np.random.multivariate_normal(np.array([1, 2]),
                                                    np.eye(2), Nsamples//2)
         dataA = np.vstack((dataAclass0, dataAclass1))
         dataBclass0 = np.copy(dataAclass0)
         dataBclass0[:, 0] = dataAclass0[:, 0] + 5
         dataBclass1 = np.copy(dataAclass1)
         dataBclass1[:, 0] = dataAclass1[:, 0] + 5
         dataB = np.vstack((dataBclass0, dataBclass1))
         correct_labels = [0] * 50 + [1] * 50

In [75]: fig, ax = plt.subplots(1, 2, figsize=(12, 5))
         ax[0].scatter(dataAclass0[:, 0], dataAclass0[:, 1], color='red')
         ax[0].scatter(dataAclass1[:, 0], dataAclass1[:, 1], color='blue')
         ax[0].grid(True)
         ax[0].axis('equal')
         ax[0].set_title('Dataset A')
         ax[1].scatter(dataBclass0[:, 0], dataBclass0[:, 1], color='red')
         ax[1].scatter(dataBclass1[:, 0], dataBclass1[:, 1], color='blue')
         ax[1].grid(True)
         ax[1].axis('equal')
         ax[1].set_title('Dataset B')
         plt.show()
```



1.1 Forward pass


```
In [43]: def z(phi, r, x):
          return 1 / (1 + np.exp(
              -np.sin(phi) * x[:, 0] - np.cos(phi) * x[:, 1] - r
          ))

In [56]: N = 50
          M = 50
          phispace = np.linspace(0, 2 * np.pi, N)
          rspace = np.linspace(-5, 5, M)

In [59]: heatA = np.zeros((M, N))
          heatB = np.zeros((M, N))
          for i, phi in enumerate(phispace):
              for j, r in enumerate(rspace):
                  heatA[i, j] = log_loss(correct_labels, z(phi, r, dataA))
                  heatB[i, j] = log_loss(correct_labels, z(phi, r, dataB))

In [82]: fig, ax = plt.subplots(1, 2, figsize=(12, 5))
          pcm = ax[0].pcolor(phispace, rspace, heatA, cmap='jet', vmin=0, vmax=5)
          ax[0].grid(True)
          ax[0].set_xlabel('$\phi$')
          ax[0].set_ylabel('$r$')
          ax[0].set_title('Dataset A')
          fig.colorbar(pcm, ax=ax[0])
          pcm = ax[1].pcolor(phispace, rspace, heatB, cmap='jet', vmin=0, vmax=5)
          ax[1].grid(True)
          ax[1].set_xlabel('$\phi$')
          ax[1].set_ylabel('$r$')
          ax[1].set_title('Dataset B')
```

```
fig.colorbar(pcm, ax=ax[1])
plt.show()
```



Observation: Works way better for dataset A in general. Also, the choice of ϕ and R is much more relevant to achieve a good result on B.

1.2 Fisher matrix

```
In [83]: def make_fisher(phi, r, x):
    mat = np.zeros((2, 2)) # [[F_phiphi, F_phir], [F_phir, F_rr]]
    term_dphi_y0 = (
        (np.cos(phi) * x[:, 0] - np.sin(phi) * x[:, 1])
        / (1 + np.exp(np.sin(phi) * x[:, 0] - np.cos(phi) * x[:, 1] + r))
    )
    term_dphi_y1 = (
        (np.cos(phi) * x[:, 0] - np.sin(phi) * x[:, 1])
        / (1 + np.exp(np.sin(phi) * x[:, 0] - np.cos(phi) * x[:, 1] + r))
        - np.cos(phi) * x[:, 0]
        + np.sin(phi) * x[:, 1]
    )
    term_dr_y0 = 1 / (1 + np.exp(
        np.sin(phi) * x[:, 0] + np.cos(phi) * x[:, 1] + r
    ))
    term_dr_y1 = 1 / (1 + np.exp(
        np.sin(phi) * x[:, 0] + np.cos(phi) * x[:, 1] + r
    )) - 1
    mat[0, 0] = np.sum(
        term_dphi_y0 * term_dphi_y0 * z(phi, r, x)
        + term_dphi_y1 * term_dphi_y1 * z(phi, r, x)
    )
```

```


mat[1, 0] = np.sum(
    term_dphi_y0 * term_dr_y0 * z(phi, r, x)
    + term_dphi_y1 * term_dr_y1 * z(phi, r, x)
)
mat[1, 1] = np.sum(
    term_dr_y0 * term_dr_y0 * z(phi, r, x)
    + term_dr_y1 * term_dr_y1 * z(phi, r, x)
)

mat[0, 1] = mat[1, 0]
return mat

fisherA = np.zeros((M, N, 2, 2))
fisherB = np.zeros((M, N, 2, 2))
for i, phi in enumerate(phispace):
    for j, r in enumerate(rspace):
        fisherA[i, j, :, :] = make_fisher(phi, r, dataA)
        fisherB[i, j, :, :] = make_fisher(phi, r, dataB)

In [95]: fig, ax = plt.subplots(2, 3, figsize=(16, 10))
for i, fisher in enumerate([fisherA, fisherB]):
    pcm = ax[i, 0].pcolor(phispace, rspace, fisher[:, :, 0, 0], cmap='jet')
    ax[i, 0].grid(True)
    ax[i, 0].set_xlabel('$\phi$')
    ax[i, 0].set_ylabel('$r$')
    fig.colorbar(pcm, ax=ax[i, 0])
    pcm = ax[i, 1].pcolor(phispace, rspace, fisher[:, :, 1, 0], cmap='jet')
    ax[i, 1].grid(True)
    ax[i, 1].set_xlabel('$\phi$')
    ax[i, 1].set_ylabel('$r$')
    fig.colorbar(pcm, ax=ax[i, 1])
    pcm = ax[i, 2].pcolor(phispace, rspace, fisher[:, :, 1, 1], cmap='jet')
    ax[i, 2].grid(True)
    ax[i, 2].set_xlabel('$\phi$')
    ax[i, 2].set_ylabel('$r$')
    fig.colorbar(pcm, ax=ax[i, 2])
ax[0, 0].set_title(r'$F_{\phi\phi}$')
ax[0, 1].set_title(r'$F_{r\phi}$')
ax[0, 2].set_title(r'$F_{rr}$')
plt.show()

```



In []: